

MindGames: Targeting Theory of Mind in Large Language Models with Dynamic Epistemic Modal Logic

Damien Sileo

Univ. Lille, Inria, CNRS

Centrale Lille, UMR 9189

CRIStAL, F-59000 Lille, France

damien.sileo@inria.fr

Antoine Lernould

Univ. Lille

CRIStAL, F-59000 Lille, France

Abstract

Theory of Mind (ToM) is a critical component of intelligence but its assessment remains the subject of heated debates. Prior research applied human ToM assessments to natural language processing models using either human-created standardized tests or rule-based templates. However, these methods primarily focus on simplistic reasoning and require further validation. Here, we leverage dynamic epistemic logic to isolate a particular component of ToM and to generate controlled problems. We also introduce new verbalization techniques to express these problems in English natural language. Our findings indicate that some language model scaling (from 70M to 6B and 350M to 174B) does not consistently yield results better than random chance. While GPT-4 demonstrates superior epistemic reasoning capabilities, there is still room for improvement. Our code and datasets are publicly available¹

1 Introduction

Theory of Mind (ToM) is the cognitive ability to attribute mental states, such as beliefs, desires, and intentions, to oneself and others, allowing individuals to understand and predict behavior based on these inferred mental states. It is an important requirement for general text understanding or artificial intelligence (Navarro et al., 2020), but claims about ToM are prone to bias from human expectations (de Waal, 2016). Kosinski (2023) recently sparked debate by showing that scaling large language models (LLMs) improves performance at standardized tests designed to measure ToM. However, these tests were widely discussed in academic research and might have leaked into the training corpora of LLM. Earlier work generated synthetic examples instead, extending the bAbi (Weston et al., 2016) framework. Nematzadeh et al. (2018) proposed a dataset of fixed templates based on the

Sally-Anne problem (Baron-Cohen et al., 1985):

Sally puts a marble in a box while Anne is with her. Sally leaves for a moment and Mary puts the marble in a basket. Where will Sally look for the marble? [ANSWER=BOX]

Le et al. (2019) deem these problems simplistic and extend them to track second-order beliefs (e.g. the belief of Sally about Anne’s beliefs).

In our study, we generate dynamic epistemic logic (DEL) problems and develop verbalizations to transform them into natural language inference problems. DEL is a branch of modal logic that can model an individual’s knowledge about particular facts or about other agents’ knowledge. DEL also enables reasoning about the impact of consecutive public announcements:

Alice and Bob have mud on their head. Their father says that at least one of them is muddy. He asks Alice and Bob if they are muddy. Do Alice and Bob know that they are muddy? [ANSWER=No] They answer that they don’t know. Do Alice and Bob now know that they are muddy? [ANSWER=Yes]

Bob would have answered YES to the first question if Alice was not muddy, so after Bob’s first answer, Alice can know that she is muddy.² DEL can formalize certain ToM problems, making it a valuable perspective for ToM assessment. The problems we create can require tracking multiple agents’ beliefs and reasoning about higher-order beliefs³. Our dataset encompasses numerous variations of the *Muddy Children* and *Drinking Logicians* problems (van Eijck, 2014). This controlled test bench offers new appreciations of language model scaling and presents the first dataset with a complexity that can challenge supervised learning models. The dataset and the scripts to generate is publicly available¹.

²The same holds if we switch Bob and Alice.

³For example, Anne’s belief about Sally’s belief about Anne’s belief about Mary’s belief.

¹[code:GitHub][data:HF-datasets]

2 Related Work

Logical Reasoning in Natural Language Processing Logic shares profound connections with NLP. Early systems were built around logic, and more recent approaches incorporate logical reasoning into neural networks (Hamilton et al., 2022; Helwe et al., 2022). Another line of research closer to ours investigates the logical capabilities of NLP models using textual datasets and labels generated with logical reasoning tools. RuleTaker (Clark et al., 2020) explores this area with propositional logic, while LogicNLI addresses first-order logic (Tian et al., 2021). Richardson and Sabharwal (2022) examine the satisfiability problem in natural language. Sileo and Moens (2022) targets probabilistic logic. Our study is the first to focus on modal logic, specifically epistemic logic, in natural language.

Theory of Mind in NLP To measure ToM capabilities of NLP models, Nematzadeh et al. (2018) created examples using Sally-Ann templates, and Le et al. (2019) added complexity to the data by incorporating second-order knowledge. Both studies framed their examples as question-answering tasks. Kosinski (2023) employed handcrafted tests to evaluate language models’ next-word prediction capabilities. Ullman (2023) showed LLM brittleness to interventions on these datasets and Ma et al. (2023) consolidated the prior datasets into a principled evaluation suite. The Social-IQA dataset (Sap et al., 2019) covers a broad spectrum of social commonsense, encompassing aspects of theory of mind and challenges like comprehending desires and emotions. Cohen (2021) investigated whether natural language inference models captured veridicality with epistemic verbs like *know* and *think*, using handcrafted patterns. This task was incorporated into the BIG-Bench framework (Srivastava et al., 2022) as the *epistemic-reasoning* task, but it targets only one shallow aspect of epistemic reasoning. Bara et al. (2021) used a Minecraft dataset for real-time belief deduction in collaborative tasks. Shapira et al. (2023b) highlighted LLM struggles in faux pas tests. Shapira et al. (2023a) conducted stress tests on LLMs’ social reasoning capabilities.

Epistemic Logic and ToM Bolander (2018) showed that the Sally-Ann problem could be modeled with epistemic logic. Van Ditmarsch and Labuschagne (2007) examined more general connections between DEL and ToM, while Dissing and Bolander (2020) demonstrated DEL’s applicability

in robotics. Van De Pol et al. (2018) explored the plausibility of epistemic logic for ToM by investigating its theoretical computational tractability.

3 Dynamic Epistemic Logic Problem Generation and Verbalization

3.1 Problem definition

Our objective is to simultaneously create dynamic epistemic logic problems and their corresponding natural language representations, with a (PREMISE, HYPOTHESIS, LABEL) format.

An epistemic logic problem can be decomposed into the following components:

Agents: A set of N individuals, each assigned a different arbitrary name.

Predicates: A set of Boolean predicates. Here, we use N predicates, one corresponding to each agent (e.g., *Alice has mud on her head*).

Observabilities: The description of each agent’s initial knowledge of the predicate values. We represent observabilities with a boolean matrix \mathcal{O} of size $N \times N$, where $\mathcal{O}_{i,j}=1$ means that agent i initially knows whether predicate j is true.

Announcements: A list of expressions (predicates or agent knowledge about predicates) that are shared to all agents. Announcements are made sequentially, and each new announcement can change what the agents know, even if it is the same announcement is repeated twice.

Hypothesis: An expression that may contain predicates and knowledge of agents about particular expressions after the announcements, given the agents, observabilities, and announcements grouped into a premise.

3.2 Setups: connecting predicate and observabilities

The concrete choice of predicates dictates the structure of observabilities. For example, the predicate “*Alice has mud on her head*” is observable by agents other than Alice, but “*Alice has mud on her hand*” could be observable by everyone. We group predicates and observabilities into what we call *setups* to generate textual descriptions. We define the following setups:

Forehead-mud setup

PREDICATE $_i$: <AGENT $_i$ >’s forehead is muddy.
 $\mathcal{O} : \text{ONES}(N) - \text{IDENTITY}(N)$

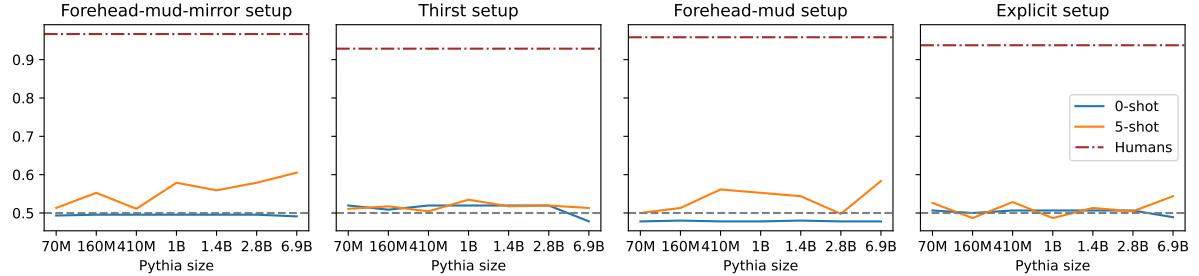


Figure 1: Accuracy of Pythia language models on MindGames setups.

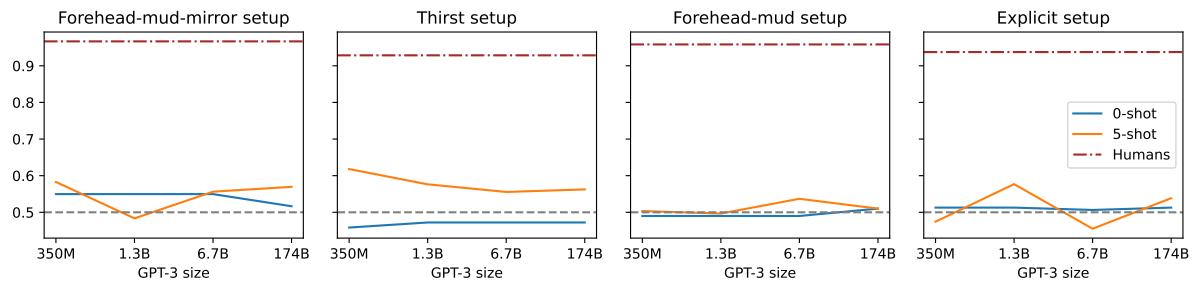


Figure 2: Accuracy of GPT-3 family (ada, cabbage, curie, davinci) language models on MindGames setups.

Forehead-mud-mirror setup

PREDICATE_i: <AGENT_i>'s forehead is muddy.
 $\mathcal{O} : \text{ONES}(N)$

OBSERVATION: There is a mirror in the room.

Thirst setup

PREDICATE_i: <AGENT_i>'s is thirsty.
 $\mathcal{O} : \text{IDENTITY}(N)$

Explicit setup

PREDICATE_i: <AGENT_i> picked a red card.

$\mathcal{O} : \text{RANDBOOL}(N, N), \mathbb{E}(\text{sum}(\mathcal{O}))=N$

OBSERVATION: Each person draws a card, face unrevealed (red or black). <<AGENT_j> card is revealed to <AGENT_i>. for all i, j where $\mathcal{O}_{i,j}=1>$

3.3 Problem verbalization

We then construct a problem for a given setup with the following natural language template:

[Premise] There are $<N>$ persons. Everyone is visible to others. <OBSERVATION> It is publicly announced that someone <PREDICATE> <[0 – N] ANNOUNCEMENTS>

[Hypothesis] <[1 – K]th ORDER BELIEF>

[0 – N] denotes uniform sampling from 0 to N. We restrict announcements to first-order beliefs. A first-order belief has the following structure: <AGENT> (can know whether | can know that | cannot know that | cannot know whether)

(<PREDICATE>|<NEGATED-PREDICATE>), e.g. Alice cannot know whether Bob is not muddy. We use can to acknowledge that an agent could theoretically infer something but fail to see it. A Kth order belief is a first-order belief about a (K–1)th order belief. We consider everyone, not everyone, and nobody as possible subjects for the setup predicates. Subjects are uniformly sampled among these quantifiers and the list of individual agents. We transform abstract problem representations into natural language and code that can be fed to a model checker to determine whether a hypothesis is entailed by the premise. We use the SMCDEL model checker (Bentham et al., 2018), an announcement logic based on the S5 (Lewis et al., 1959) modal logic. This implementation is the most cited publicly available epistemic logic as of April 2023. We discard examples where the premise contains a contradiction⁴. To generate diverse and gender-balanced random English surnames, we use CensusName⁵ (Qian et al., 2022).

4 Experiments

4.1 Problem generation parameters

We randomly sample $N \in \{2, 3, 4\}$ agents, as we observed that problems were sufficiently challeng-

⁴We identify contradictions by examining whether an unused predicate is entailed or not by the premise.

⁵<https://pypi.org/project/censusname/>